

Homework 1

In [1]:

```
%load_ext nb_black
```

In [2]:

```
import math
import numpy as np
import scipy as sc
from scipy import ndimage
from skimage import filters
from PIL import Image
from IPython.display import display

import plotly.express as px
```

Section 1: Coding Assignment

Problem 1

In [3]:

```
image = Image.open("hw1_data/Seattle.jpg") # .convert("L")
image = np.asarray(image)
print(image.shape)
px.imshow(image, color_continuous_scale="gray")
```

(600, 800, 3)



Gaussian filter is defined as:

$$G(\mathbf{x}; \sigma^2) = \frac{1}{2\pi\sigma^2} \exp -\frac{\mathbf{x}^T \mathbf{x}}{2\sigma^2}$$

In [4]:

```
# Trick for colors
# np.stack([np.arange(16).reshape(4,4)]*3, axis = -1).shape

# # First lets write a convolve function
def convolve(image, filter):
    filter_size = filter.shape[0]

    padimage = imagepadding(image, filter_size)
    rowsize = image.shape[0]
    colszie = image.shape[1]

    im_filtered = np.zeros_like(padimage, dtype=np.float32)

    for i in range(rowsize): # rows
        for j in range(colszie): # columns
            submat = padimage[i : i + filter_size, j : j + filter_size]
            im_filtered[i, j] = np.sum(np.multiply(submat, filter))
    return im_filtered

# Image Padding function
def imagepadding(image, sizeofpad):
    sizeofpad = sizeofpad // 2
    rowsize = image.shape[0]
    colszie = image.shape[1]

    # Multiply by two since padding is double sided
    padimage = np.zeros(shape=(rowsize + 2 * sizeofpad, colszie + 2 * sizeofpad))

    # Plant the image in the middle of the new matrix, equidistant from the edges
    padimage[sizeofpad:-sizeofpad, sizeofpad:-sizeofpad] = image
    return padimage
```

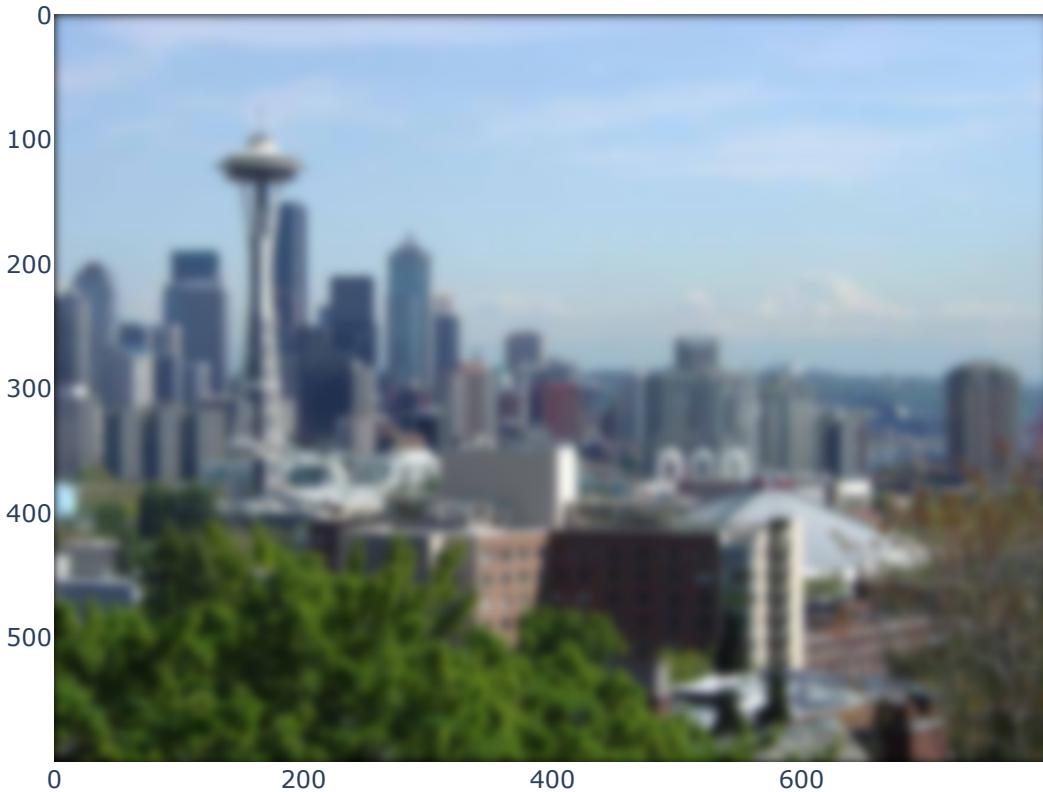
In [5]:

```
def GaussianBlurImage(image, sigma):
    # Builds the filter
    filter_size = 2 * int(4 * sigma + 0.5) + 1
    gaussian_filter = np.zeros((filter_size, filter_size), np.float32)
    for i in range(filter_size):
        for j in range(filter_size):
            x = i - filter_size // 2
            y = j - filter_size // 2
            gaussian_filter[i, j] = (
                1
                / (2 * np.pi * sigma ** 2)
                * np.exp(-(x ** 2 + y ** 2) / (2 * sigma ** 2)))
    # Zero matrix
    im_filtered = np.zeros(
        shape=(
            image.shape[0] + gaussian_filter.shape[0] - 1,
            image.shape[1] + gaussian_filter.shape[0] - 1,
            3,
```

```
)  
  
    # Filter for each RGB  
    for c in range(image.shape[2]):  
        im_filtered[:, :, c] = convolve(image[:, :, c], gaussian_filter)  
  
    # RGB version  
    im_filtered = im_filtered[: image.shape[0], : image.shape[1], :]  
]  
  
print("The final resolution is ", im_filtered.shape)  
fig = px.imshow(np.clip(im_filtered, 0, 255).astype(np.uint8))  
fig.write_image("1.png")  
display(fig)  
return im_filtered
```

```
gb = GaussianBlurImage(image, 4)
```

```
The final resolution is (600, 800, 3)
```



In [6]:

```
# im_filtered_skimage = filters.gaussian(  
#     image, sigma=4, mode="constant", cval=0.0, multichannel=True, preserve_range=True  
# )  
# print(im_filtered_skimage.shape)  
# px.imshow(im_filtered_skimage)
```

Problem 2

In [7]:

```
def SeparableGaussianBlurImage(image, sigma):
    # Create filter
    filter_size = 2 * int(4 * sigma + 0.5) + 1
    gaussian_filter = np.fromfunction(
        lambda x: math.e
        ** ((-1 * (x - (filter_size - 1) / 2) ** 2) / (2 * sigma ** 2)),
        (filter_size,),
    )
    gaussian_filter = gaussian_filter / np.sum(
        gaussian_filter
    ) # Normalizing for sum to 1

    # Pad Image for rgb
    padimage = np.zeros(
        (image.shape[0] + filter_size - 1, image.shape[1] + filter_size - 1, 3)
    )
    for rgb in range(3):
        padimage[:, :, rgb] = imagepadding(image[:, :, rgb], filter_size)
    # Create empty matrix
    im.blur = np.zeros(
        (padimage.shape[0] - filter_size + 1, padimage.shape[1] - filter_size + 1, 3)
    )

    for c in range(3):
        img = padimage[:, :, c].astype(np.uint8)
        rowblur = np.zeros((img.shape[0], img.shape[1] - filter_size + 1))
        for i, val in enumerate(gaussian_filter):
            rowblur += val * img[:, i : img.shape[1] - filter_size + i + 1]
        colblur = np.zeros((rowblur.shape[0] - filter_size + 1, rowblur.shape[1]))
        for i, val in enumerate(gaussian_filter):
            colblur += val * rowblur[i : img.shape[0] - filter_size + i + 1]
        im.blur[:, :, c] = colblur

    fig = px.imshow(im.blur)
    fig.write_image("2.png")
    display(fig)

    return im.blur

sep = SeparableGaussianBlurImage(image, 4.0)
```





Problem 3

In [8]:

```
yosemite = Image.open("hw1_data/Yosemite.png") # .convert("L")
yosemite = np.asarray(yosemite)
px.imshow(yosemite, color_continuous_scale="gray")
```



In [9]:

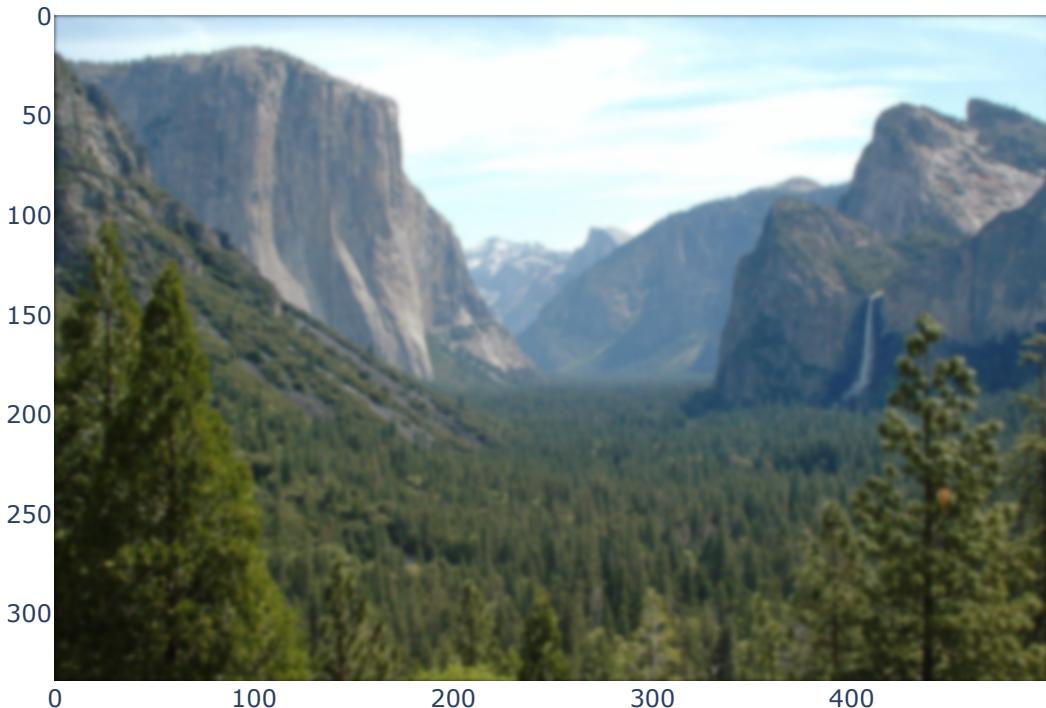
```
def SharpenImage(image, sigma, alpha):
    image = image.astype(np.uint8)
    blur = GaussianBlurImage(image, sigma)
    edges = blur - image

    # Replicating Matlabs functionality for uint8()
    edges[edges < 0] = 0
    edges[edges > 255] = 255
```

```
sharp = image.astype(np.uint8) - alpha * (edges) # im_filtered
fig = px.imshow(sharp)
fig.write_image("4.png")
return fig
```

```
SharpenImage(yosemite, 1, 5)
```

The final resolution is (334, 500, 3)





Problem 4

In [10]:

```
ladybug = Image.open("hw1_data/LadyBug.jpg").convert("L")
ladybug = np.asarray(ladybug)
px.imshow(ladybug, color_continuous_scale="gray")
```



In [24]:

```
def SobelImage(image):
    horiz = np.array([[1, 0, -1], [2, 0, -2], [1, 0, -1]])
    vert = np.array([[1, 2, 1], [0, 0, 0], [-1, -2, -1]])

    gx = convolve(image, horiz)
    gy = convolve(image, vert)

    # Convert 0 to machine epsilon
    gx[gx == 0] = np.finfo(float).eps # 2.22044604925e-16
    # gy[gy == 0] = np.finfo(float).eps
```

```

mag = np.sqrt(gx ** 2 + gy ** 2)
orient = np.arctan(gy / gx)

# Replicating Matlabs functionality for uint8()
mag[mag < 0] = 0
mag[mag > 255] = 255
px.imshow(mag, color_continuous_scale="gray")

#      orient[orient < 0] = 0
#      orient[orient > 255] = 255

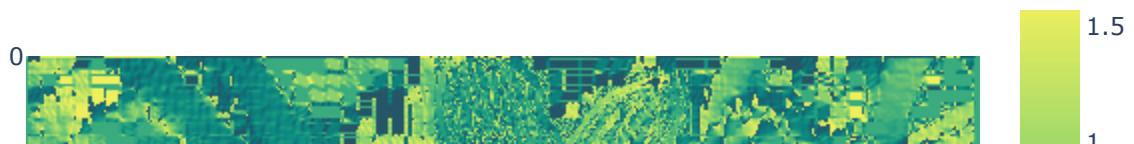
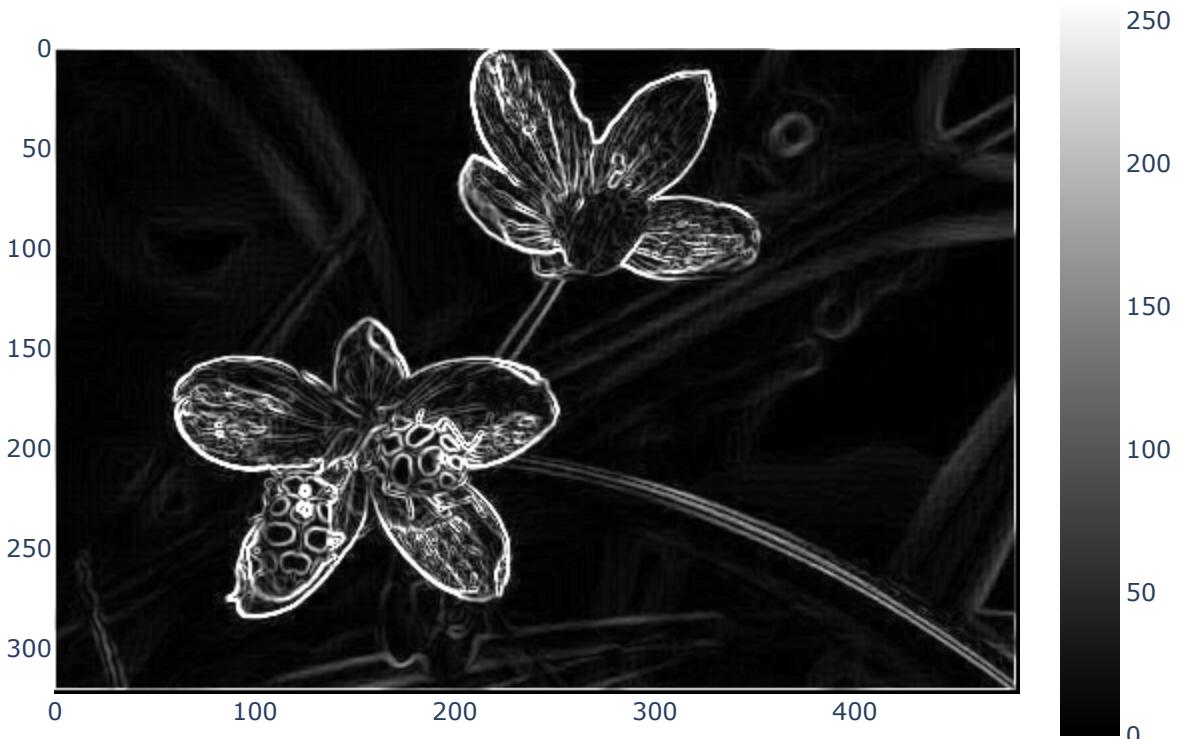
fig1 = px.imshow(mag, color_continuous_scale="gray")
fig2 = px.imshow(orient, color_continuous_scale="aggrnyl")

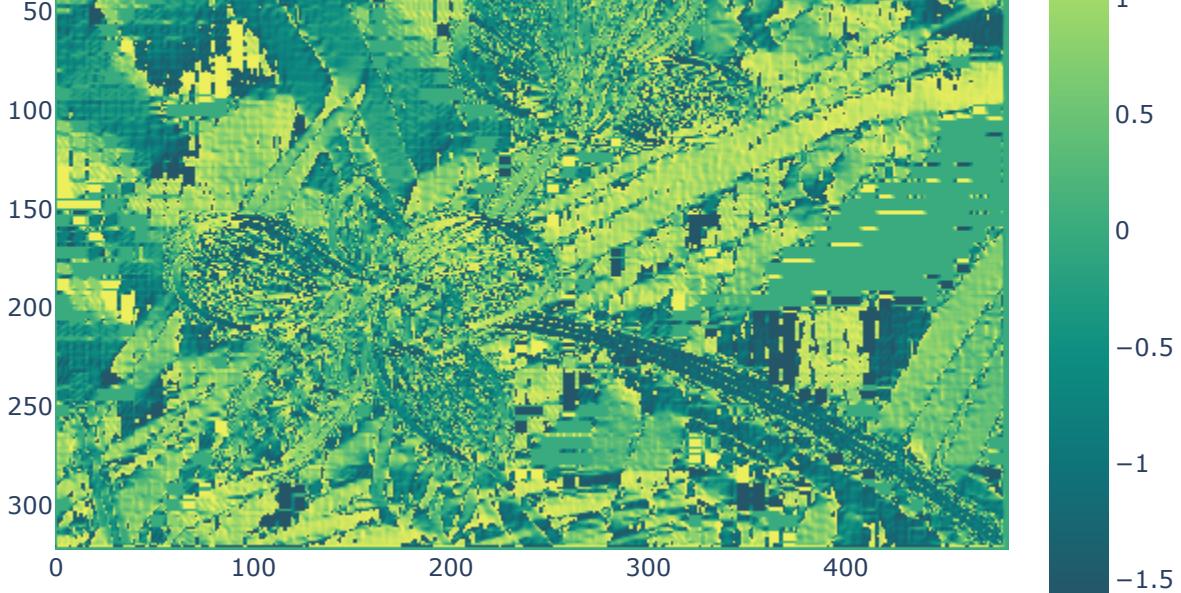
fig1.write_image("5a.png")
fig2.write_image("5b.png")

display(fig1)
display(fig2)
return (mag, orient)

```

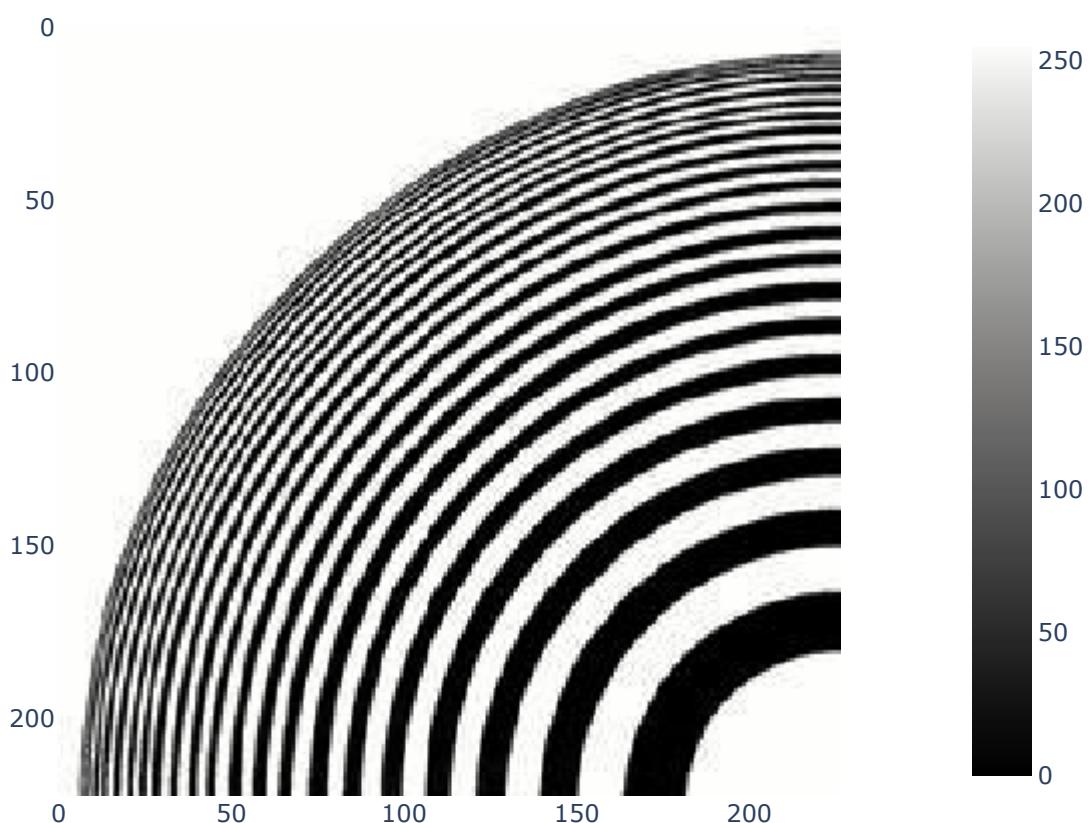
```
magnitude, orientation = SobelImage(ladybug)
```





Problem 5

```
In [12]:  
msmall = Image.open("hw1_data/Moire_small.jpg").convert("L")  
msmall = np.asarray(msmall)  
px.imshow(msmall, color_continuous_scale="gray")
```



In [151...]

```

def BilinearInterpolation(image, x, y, scalefac=4, method="NN", imagewrite=True):
    size = np.shape(image)
    im_scaled = np.zeros(shape=(size[0] * scalefac, size[1] * scalefac), dtype=np.uint8)
    if method == "NN":
        for i in range(scalefac * size[0]):
            for j in range(scalefac * size[1]):
                x1 = int(np.floor(i / scalefac))
                y1 = int(np.floor(j / scalefac))
                im_scaled[i, j] = image[x1, y1]

        fig = px.imshow(im_scaled, color_continuous_scale="gray")
        display(fig)
        fig.write_image("6a.png")
    else: # Bilinear
        for i in range(
            size[0] * scalefac - 3
        ): # I am not quite sure why minus 3 works yet minus 2 doesn't
            # print(i)
            x1 = int(math.floor(i / scalefac)) # x
            xp1 = int(math.ceil(i / scalefac)) # x plus 1
            a = i / 4 % 1

            for j in range(size[1] * scalefac - 3):
                y1 = int(math.floor(j / scalefac)) # y
                yp1 = int(math.ceil(j / scalefac)) # y plus 1
                b = j / 4 % 1

                # Corners
                topleft = image[x1, y1]
                topright = image[xp1, y1]
                botleft = image[x1, yp1]
                botright = image[xp1, yp1]

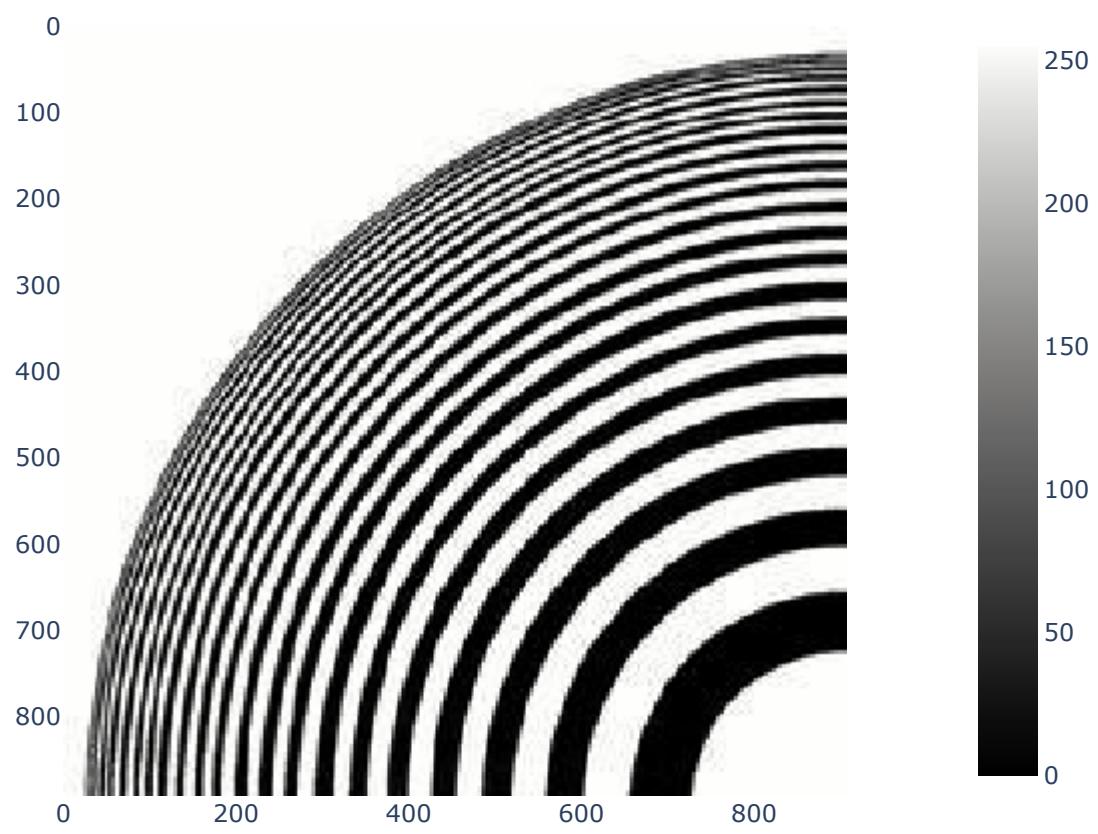
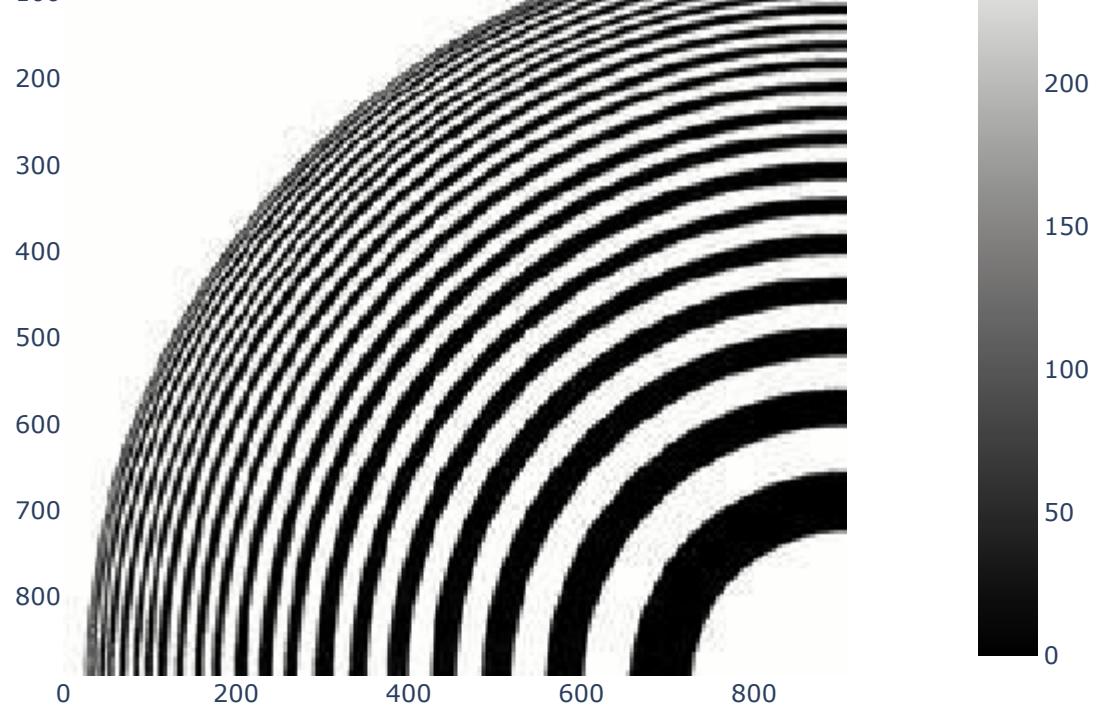
                newval = (
                    ((1 - a) * (1 - b) * topleft)
                    + (a * (1 - b) * topright)
                    + (b * (1 - a) * botleft)
                    + (a * b * botright)
                )
                im_scaled[i, j] = newval

    if imagewrite == True:
        fig = px.imshow(im_scaled, color_continuous_scale="gray")
        display(fig)
        fig.write_image("6b.png")

    #     print("Original resolution", image.shape)
    #     print("New Resolution", im_scaled.shape)
    return (x, y)

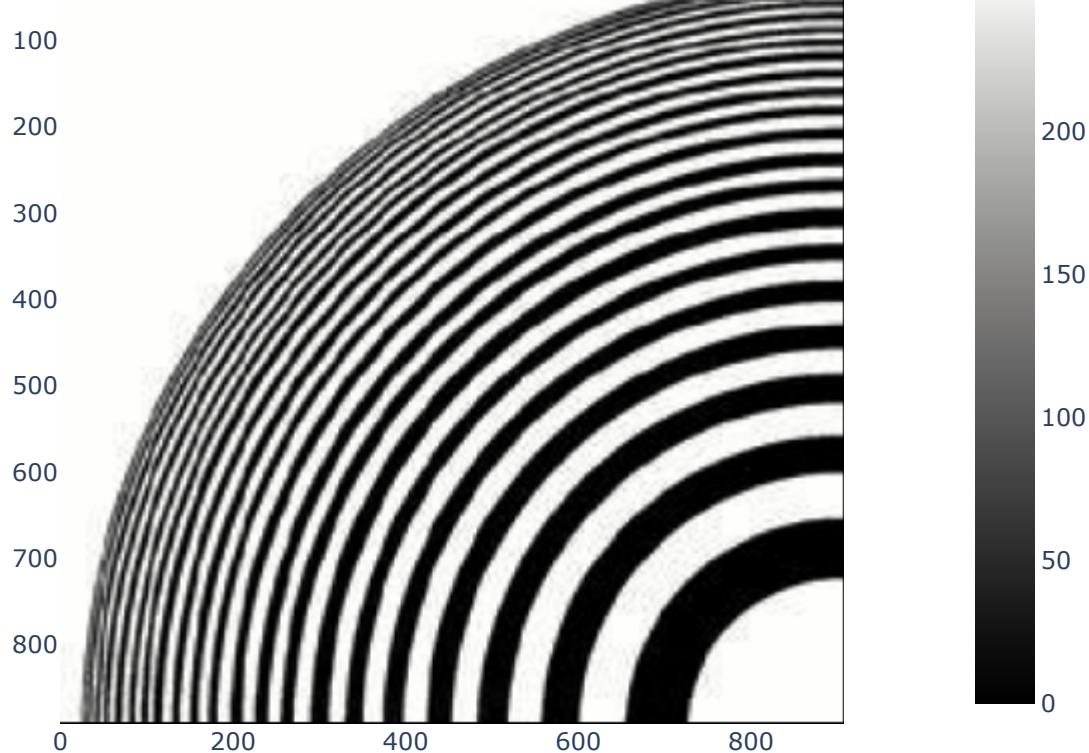
bilinNN = BilinearInterpolation(msmall, 465, 518, method="NN")
print(bilinNN)
bilinBL = BilinearInterpolation(msmall, 465, 518, method="Bilinear")
print(bilinBL)

```



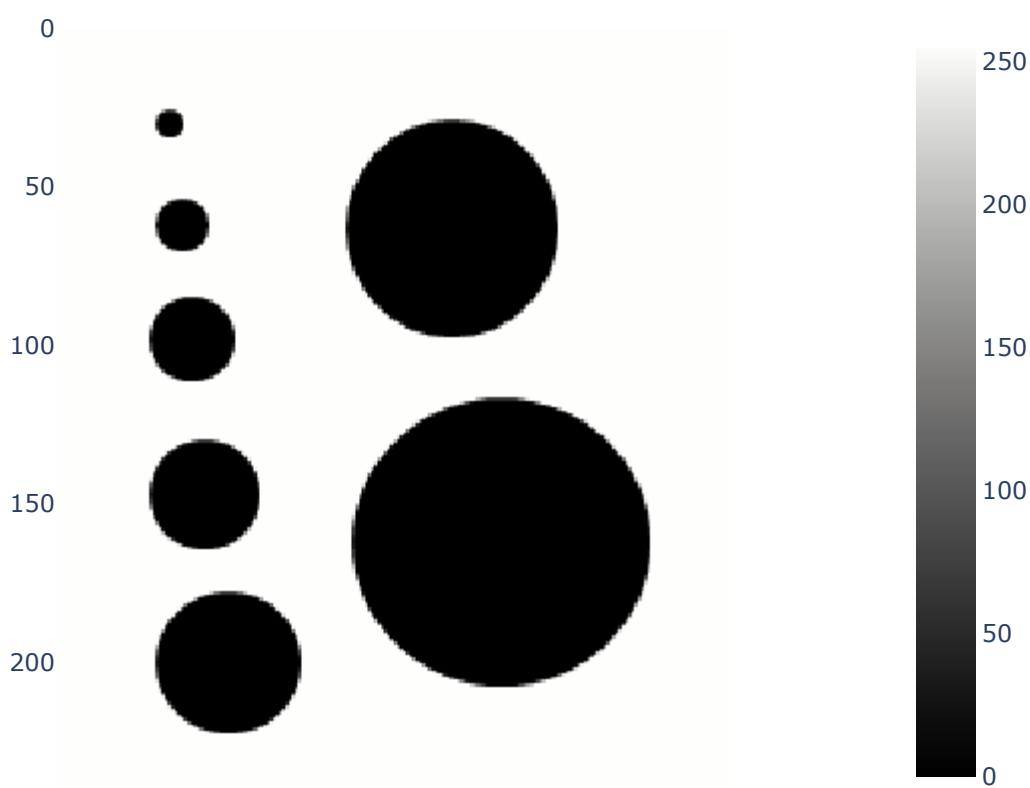
(465, 518)





Question 6

```
In [14]: circles = Image.open("hw1_data/Circle.png").convert("L")
circles = np.asarray(circles)
px.imshow(circles, color_continuous_scale="gray")
```



In [162...]

```

def FindPeaksImage(image, thres):
    mag, theta = SobelImage(image)
    # print(orient)

    for i in range(
        mag.shape[0] - 3
    ): # minus three since we do not want to go over bounds of matrix
        for j in range(mag.shape[1] - 3):
            # e0, e1 = BilinearInterpolation(
            #     image, i + 1, j + 1, method="Bilinear", imagewrite=False
            # )

            submag = mag[i : i + 3, j : j + 3]
            subtheta = theta[i : i + 3, j : j + 3]

            # find which cells are in same direction as center
            centermag = submag[1, 1]
            centertheta = subtheta[1, 1]

            neighbors = np.argwhere(subtheta == centertheta)
            # print(neighbors)
            # Remove the center one (redundant neighbor to itself)
            mid = [1, 1]
            neighbors = [i for i in neighbors if not np.any(i == mid)]

            if len(neighbors) == 2:
                n1 = neighbors[0] # first neighbor
                n2 = neighbors[1] # second neighbor

                # now that same direction cells are found, look at intensity
                # if less than center magnitude and threshold, turn to zero
                # here we compare the first neighbor
                if (submag[n1[0], n1[1]] < thres) & (submag[n1[0], n1[1]] < centermag):
                    submag[n1[0], n1[1]] == 0
                if (submag[n1[0], n1[1]] > thres) & (submag[n1[0], n1[1]] > centermag):
                    # change the center to zero since it is not as intense as neighbor
                    submag[1, 1] == 0
                    submag[n1[0], n1[1]] == 255
                # Now second neighbor
                if (submag[n2[0], n2[1]] < thres) & (submag[n2[0], n2[1]] < centermag):
                    submag[n2[0], n2[1]] == 0
                if (submag[n2[0], n2[1]] > thres) & (submag[n2[0], n2[1]] > centermag):
                    submag[1, 1] == 0
                    submag[n2[0], n2[1]] == 255

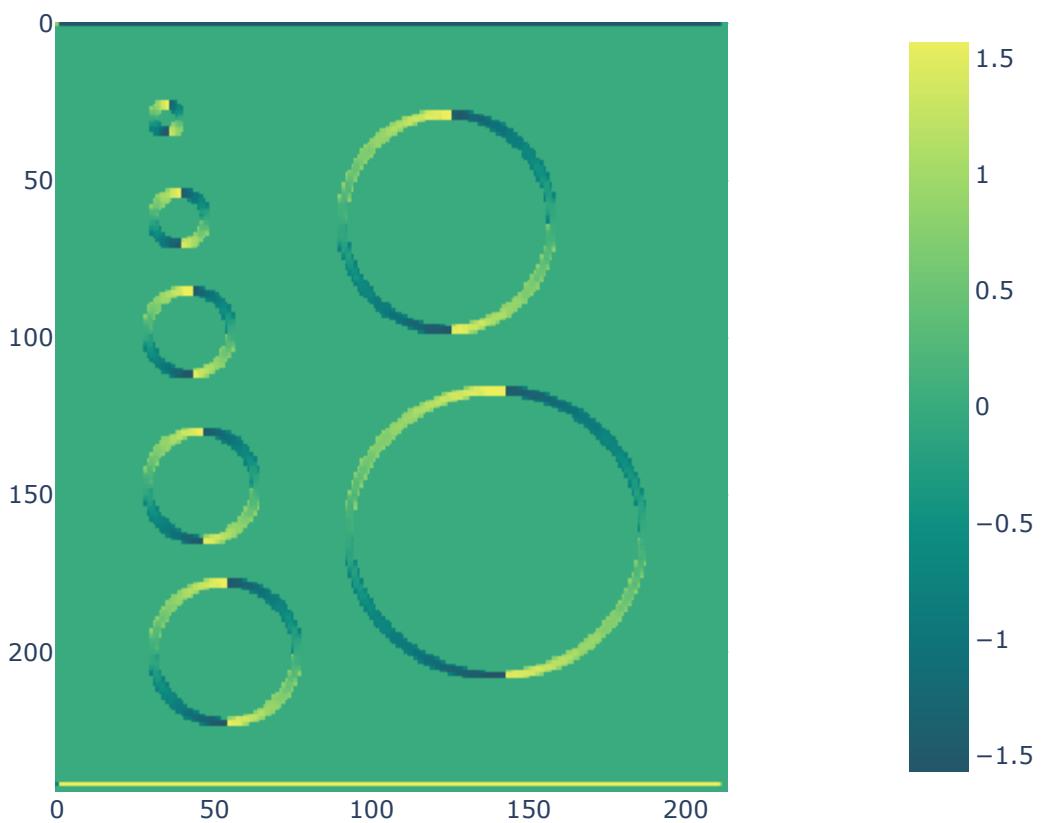
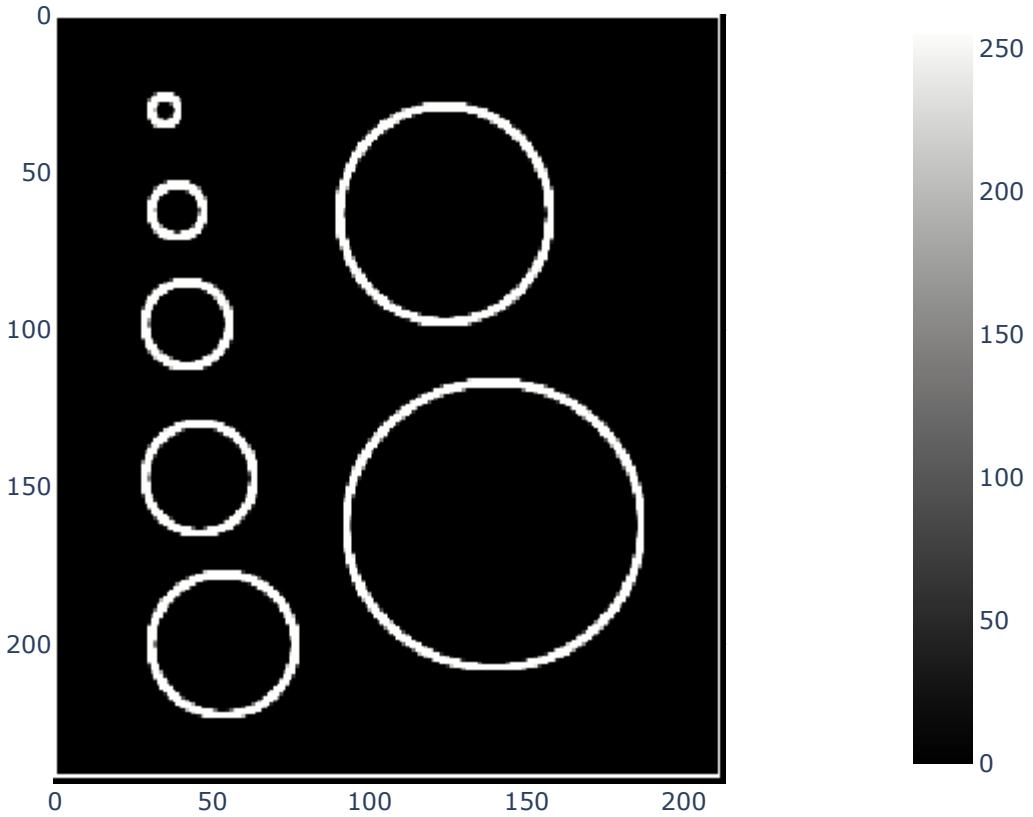
            mag[i : i + 3, j : j + 3] = submag

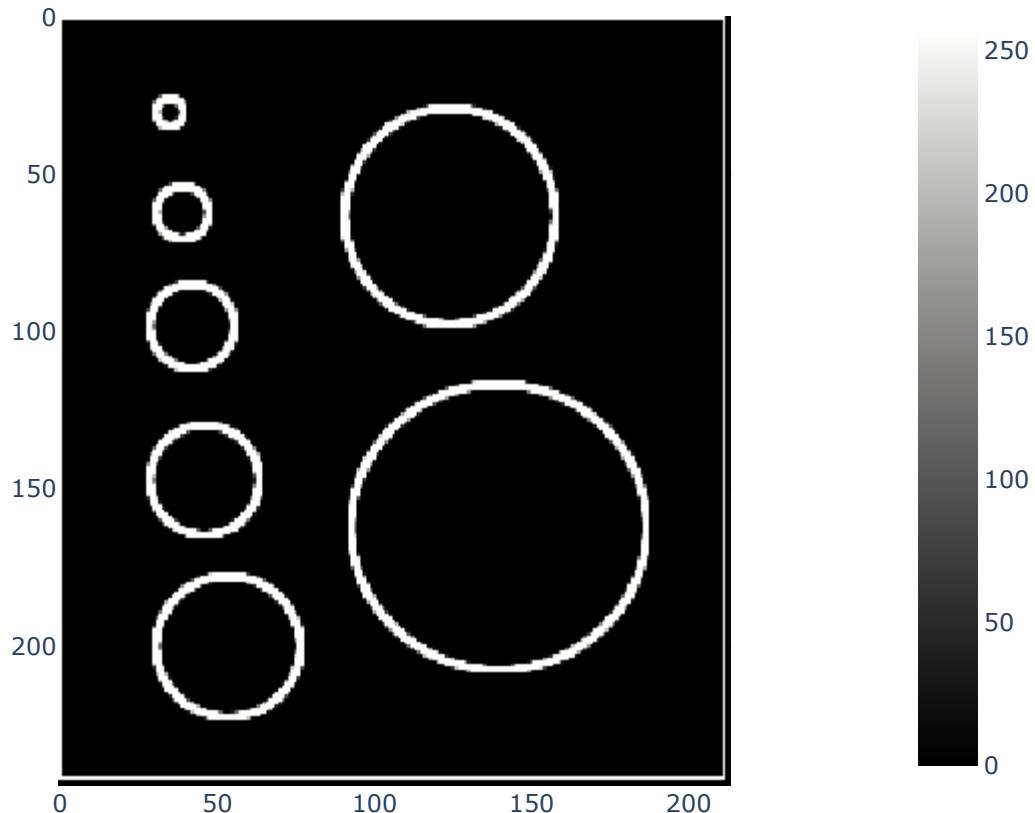
fig = px.imshow(mag)
fig.write_image("7.png")

return mag

```

px.imshow(FindPeaksImage(circles, 40), color_continuous_scale="gray")





Section 2: Written Assignment

Question 1

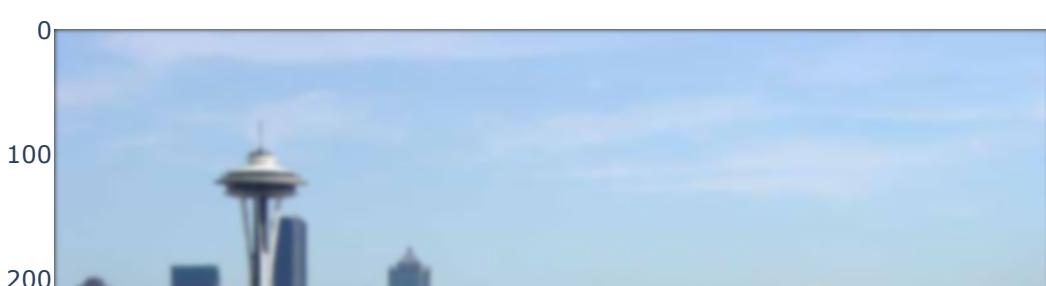
In [16]:

```
import time
```

In [17]:

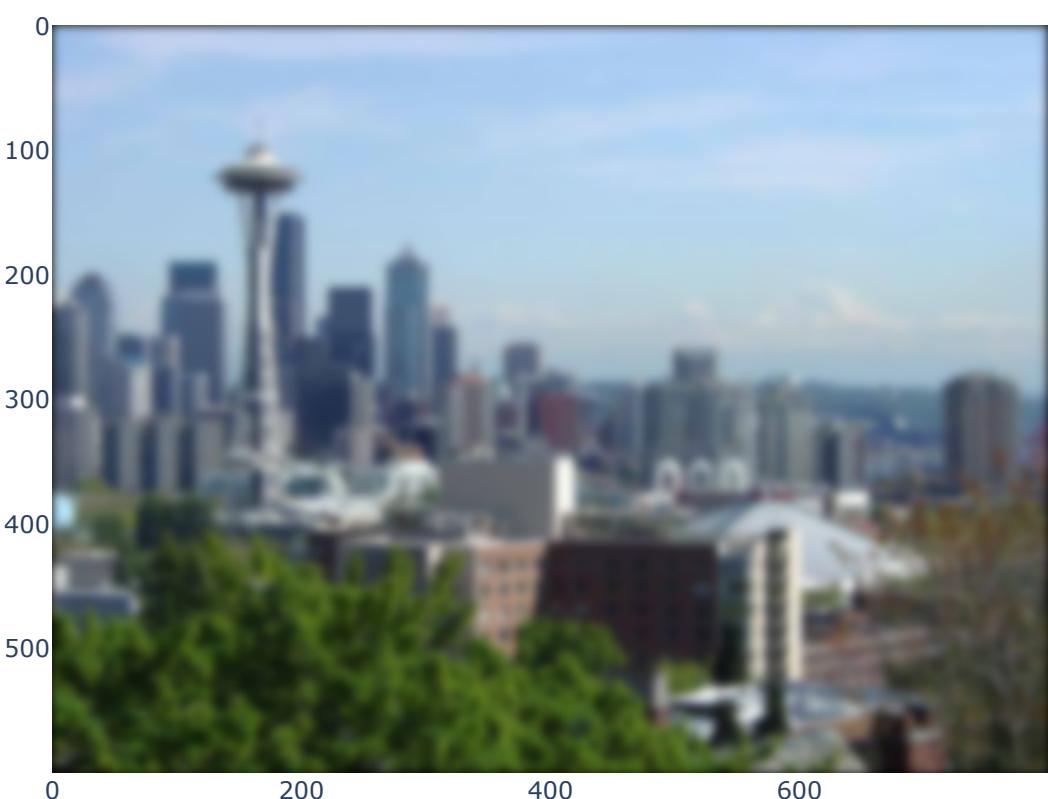
```
start = time.time()
GaussianBlurImage(image, 2)
print("--- %s seconds ---" % (time.time() - start))
start = time.time()
GaussianBlurImage(image, 4)
print("--- %s seconds ---" % (time.time() - start))
start = time.time()
GaussianBlurImage(image, 8)
print("--- %s seconds ---" % (time.time() - start))
```

The final resolution is (600, 800, 3)



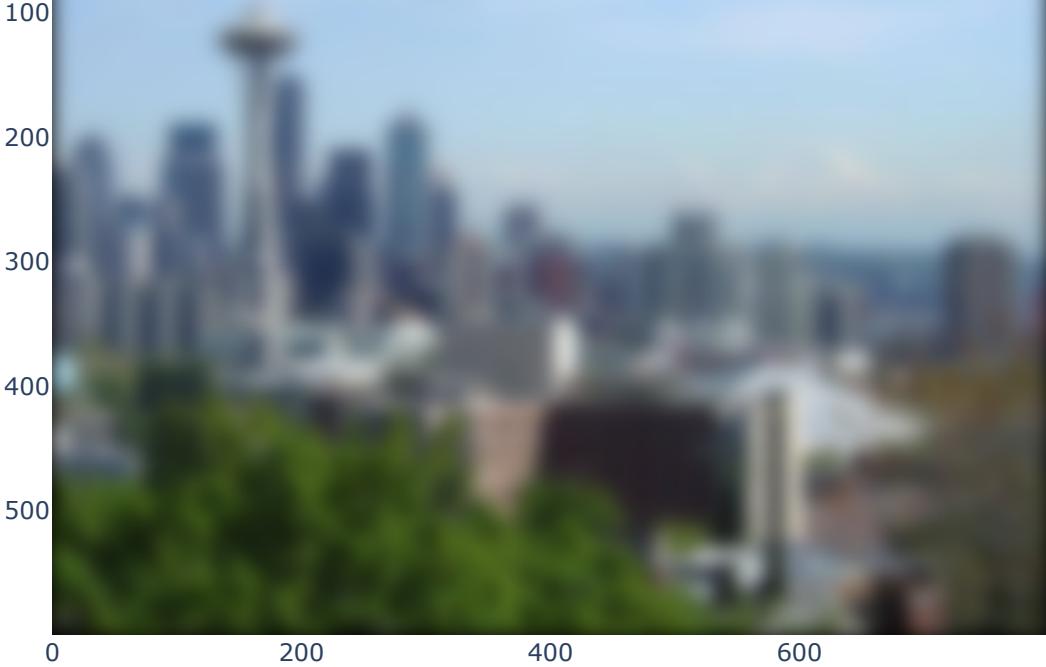


--- 7.835134029388428 seconds ---
The final resolution is (600, 800, 3)



--- 9.607639074325562 seconds ---
The final resolution is (600, 800, 3)





--- 17.2693293094635 seconds ---

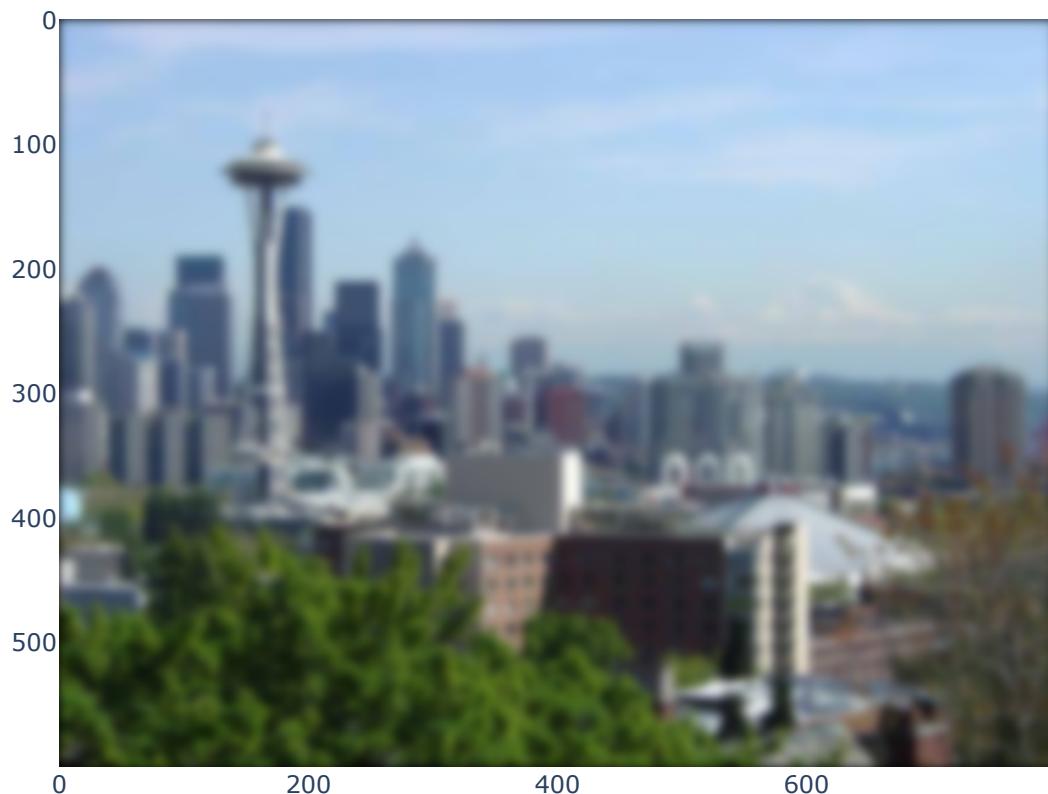
In [18]:

```
start = time.time()
SeparableGaussianBlurImage(image, 2)
print("--- %s seconds ---" % (time.time() - start))
start = time.time()
SeparableGaussianBlurImage(image, 4)
print("--- %s seconds ---" % (time.time() - start))
start = time.time()
SeparableGaussianBlurImage(image, 8)
print("--- %s seconds ---" % (time.time() - start))
```

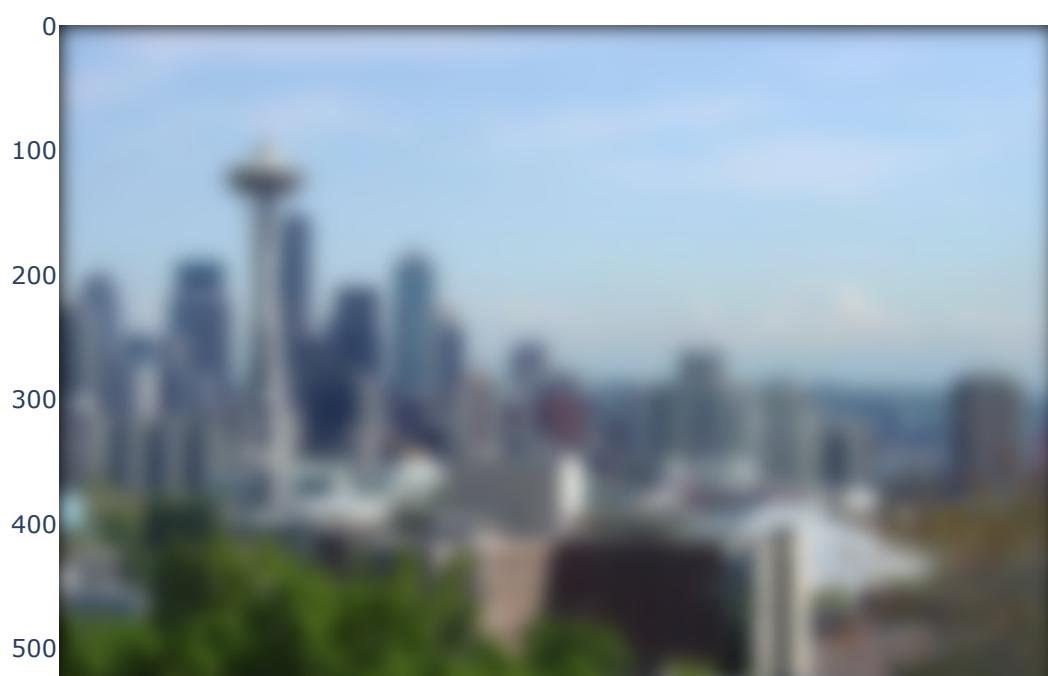




--- 0.6198980808258057 seconds ---



--- 0.7106447219848633 seconds ---





0 200 400 600

--- 0.8921654224395752 seconds ---

For a $\sigma = 32$ I believe it would take over 30 seconds for GaussianBlurImage() to perform the task and over 2 seconds for SeparableGaussianBlur() to finish. We can see from above that an increase in σ increases the run time of the blurring.

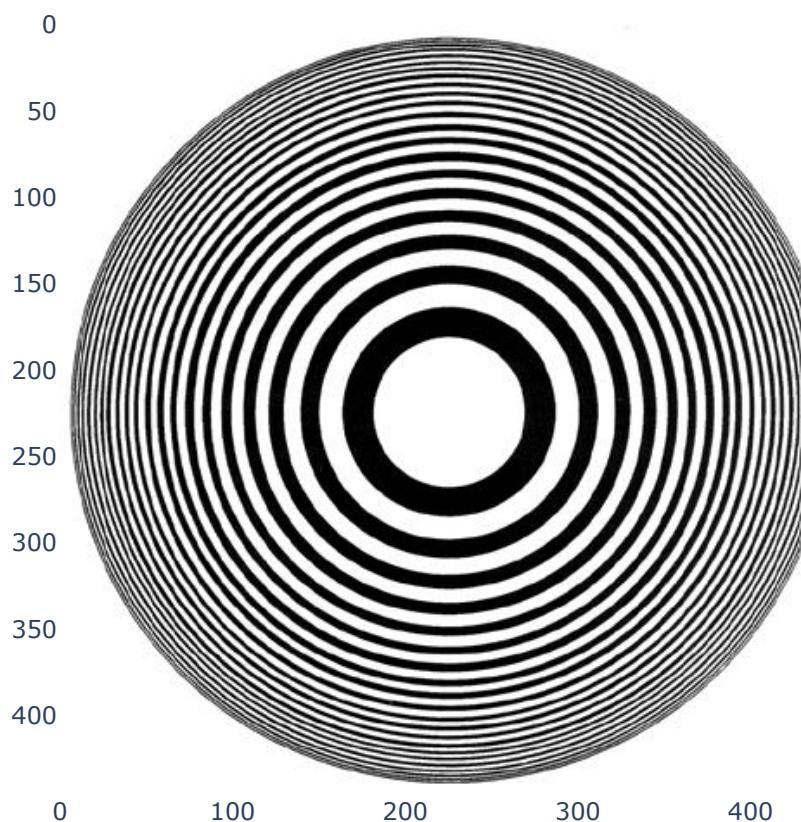
Question 2

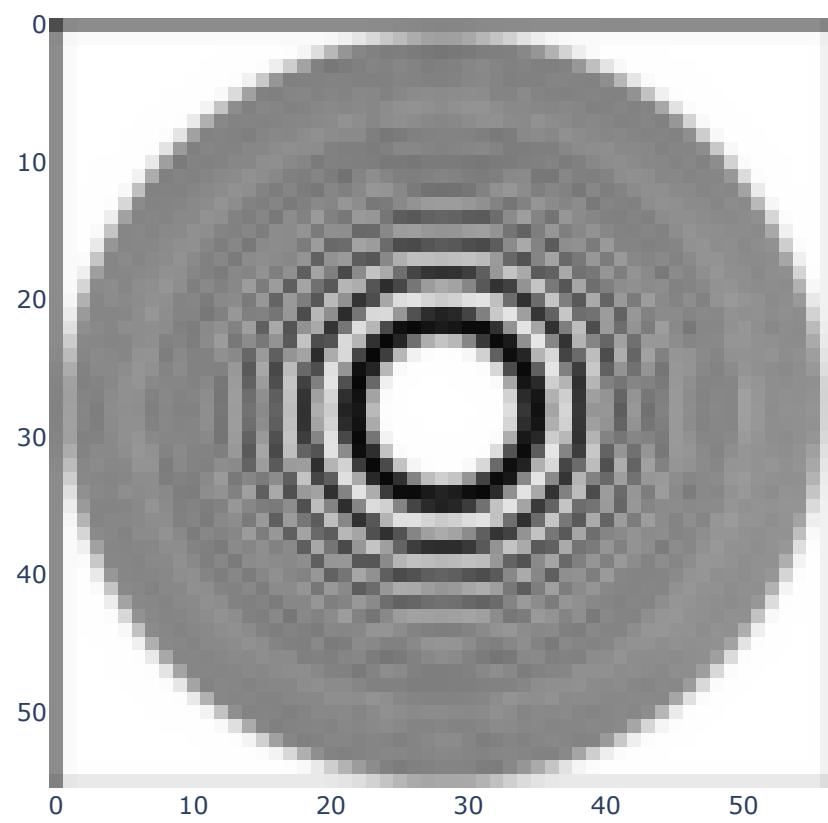
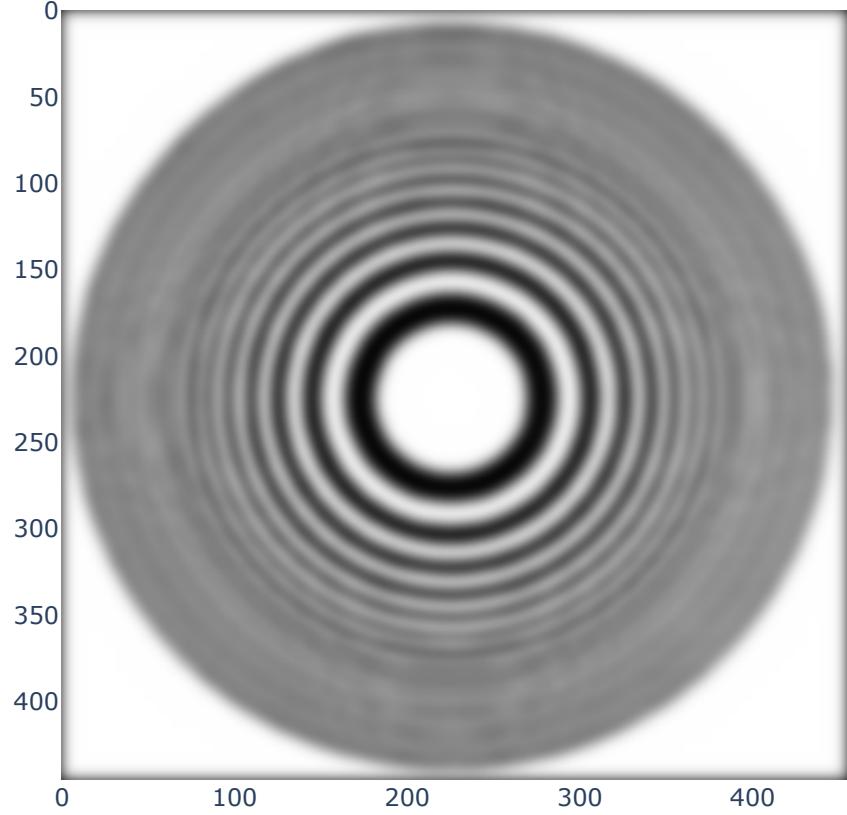
In [180...]

```
moire = Image.open("hw1_data/Moire.jpg")    # .convert("L")
moire = np.asarray(moire)
display(px.imshow(moire, color_continuous_scale="gray"))
moire = SeparableGaussianBlurImage(moire, 4)

for i in range(3):    # Half size 3 times down sampling
    moire = moire[0::2, 0::2]

display(px.imshow(moire, color_continuous_scale="gray"))
```



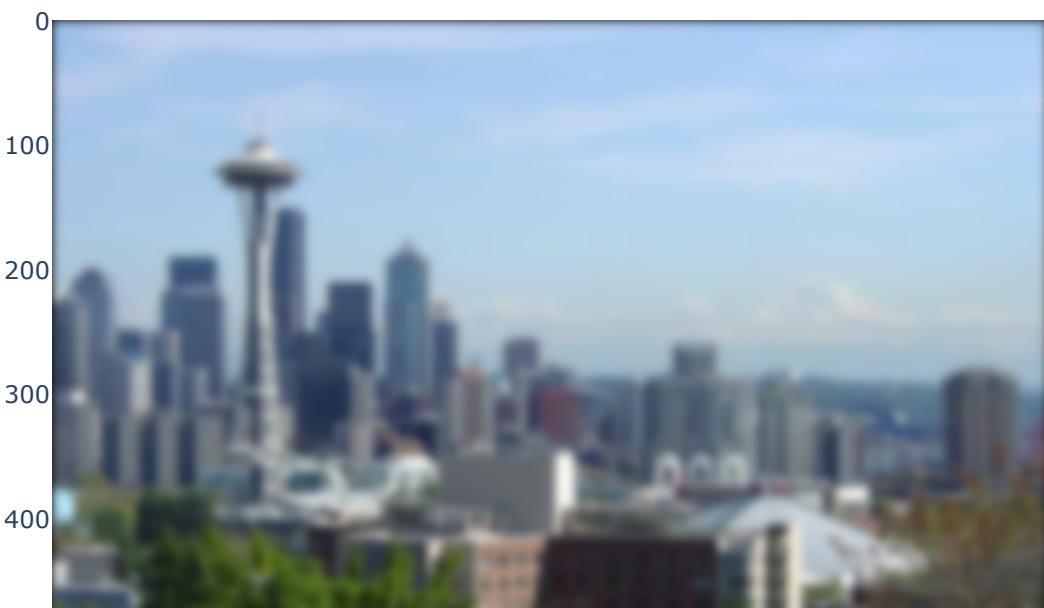


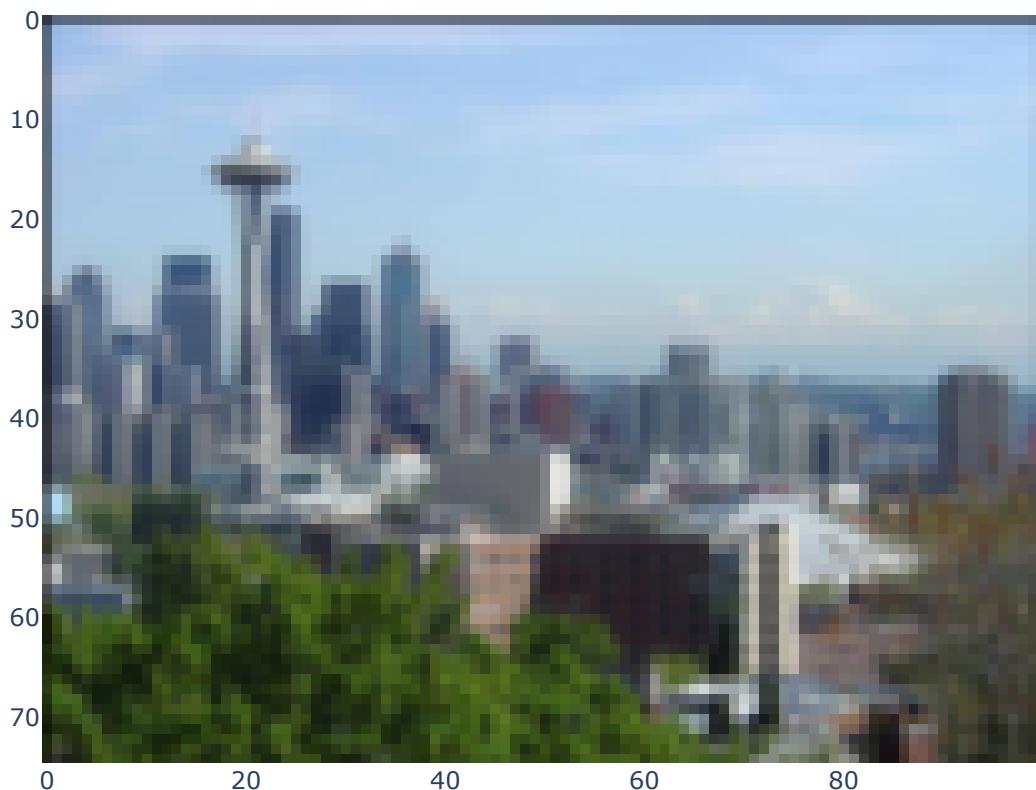
In [182]:

```
seattle = Image.open("hw1_data/seattle.jpg") # .convert("L")
seattle = np.asarray(seattle)
display(px.imshow(seattle, color_continuous_scale="gray"))
seattle = SeparableGaussianBlurImage(seattle, 4)

for i in range(3): # Half size 3 times down sampling
    seattle = seattle[0::2, 0::2]

display(px.imshow(seattle, color_continuous_scale="gray"))
```





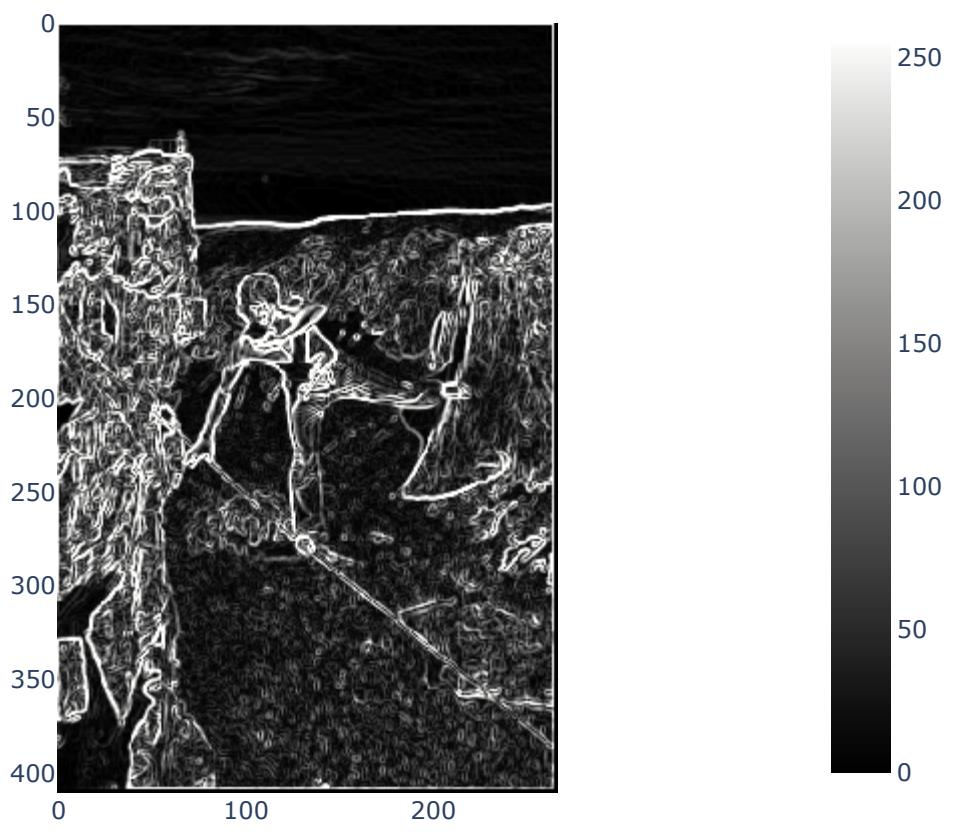
The best amount of blur for the Moire image seems to be with a sigma of 4, however, this value does not perform the most ideal output. In this question, I am assuming best is that we can blur an image with a downsampled version of the image and maintain a optimum amount of information such that the blurring looks very close to blurring the original image.

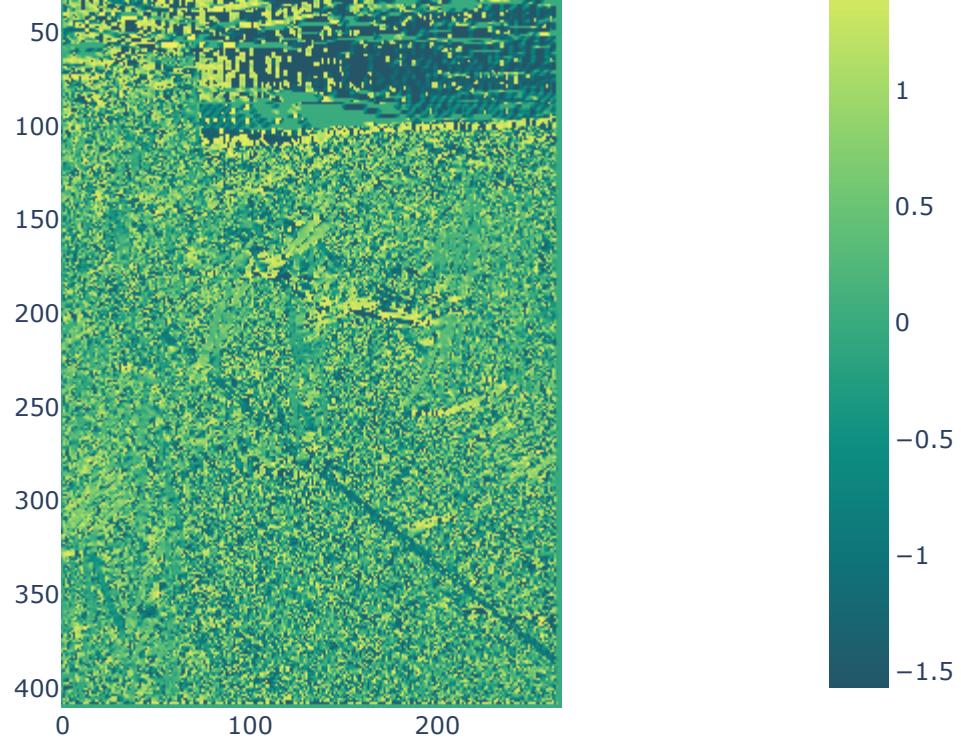
Question 3

In [26]:

```
tightrope = Image.open("hw1_data/TightRope.png").convert("L")
tightrope = np.asarray(tightrope)
display(px.imshow(tightrope, color_continuous_scale="gray"))
x = SobelImage(tightrope)
```







From the tight rope image, we can see that an obvious edge would be the cliff at the top left, however, the Sobel filters do not give this edge a lot of strength in the orientation. Another not so well seen edge to the filter is the side of the man's left leg.

Question 4

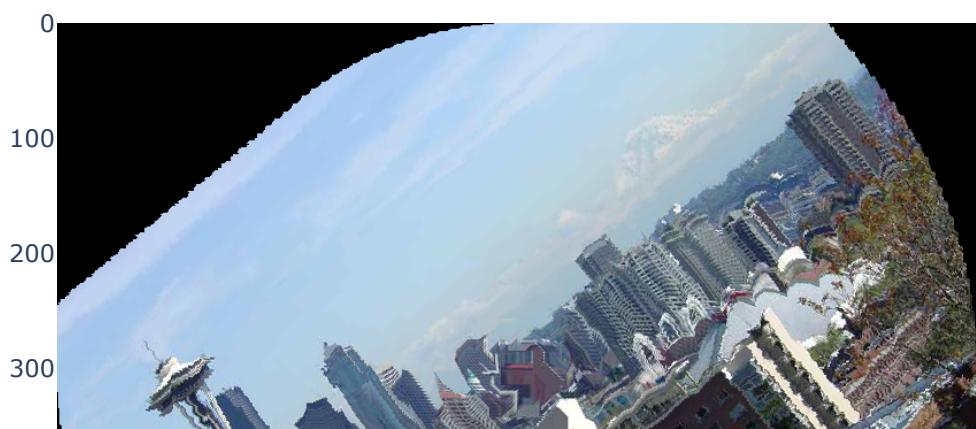
In [160...]

```
seattle = Image.open("hw1_data/seattle.jpg")
for i in range(20):
    seattle = seattle.rotate(2)

display(px.imshow(seattle, title="Rotated by 2 Degrees 20 times"))

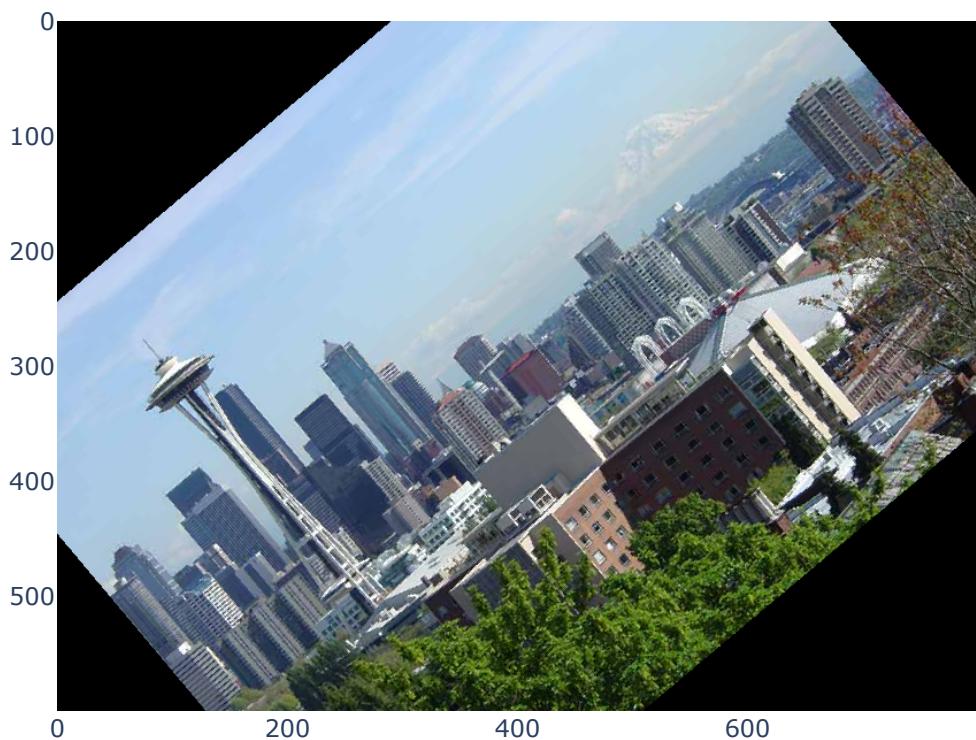
seattle = Image.open("hw1_data/seattle.jpg")
display(px.imshow(seattle.rotate(40), title="Rotated 40 Degrees"))
```

Rotated by 2 Degrees 20 times





Rotated 40 Degrees



As we can see here, rotating an image 20 times by 2 degrees yields a different result than rotating an image 40 degrees. Performing so many small rotations begins to misplace pixels each rotation causing a warping effect. A single rotation maintains the quality of the original image without making any visible warping.